PRISM.

See what survives the scrub

A privacy engineering workbench
for enterprise data-sharing teams.

The names are gone.
The retained detail still matters.

A research team wants to share a table with an internal or external collaborator.
Our fictional oncology demo shows how retained age, geography and weight can distinguish records.

01

Source table

Review identifiers,
known attributes and units.

02

Candidate release

Remove identifiers.
Generalize retained detail.

03

Study questions

Measure exposure and
changes to selected answers.

PRISM ties disclosure and utility review to the team's declared purpose.

Six connected views.
One candidate under review.

01	**Workbench**	Inspect the source, transformed table and scientific comparisons.
02	**Knowledge Map**	See which declared attribute combinations distinguish records.
03	**Research Intent**	Choose study questions, tolerances and minimum cohort sizes.
04	**Pareto Explorer**	Compare actual recipes across measured exposure and utility loss.
05	**Attack Lab**	Explore fictional linkage, composition and correlation counterexamples.
06	**Report**	Export evidence, the recipe and a reviewed protected CSV.

Source and candidate remain distinct. Research intent defines the comparison.

55 fewer unique correct matches

One change: remove ZIP from the protected candidate.

Declared exact-match attack

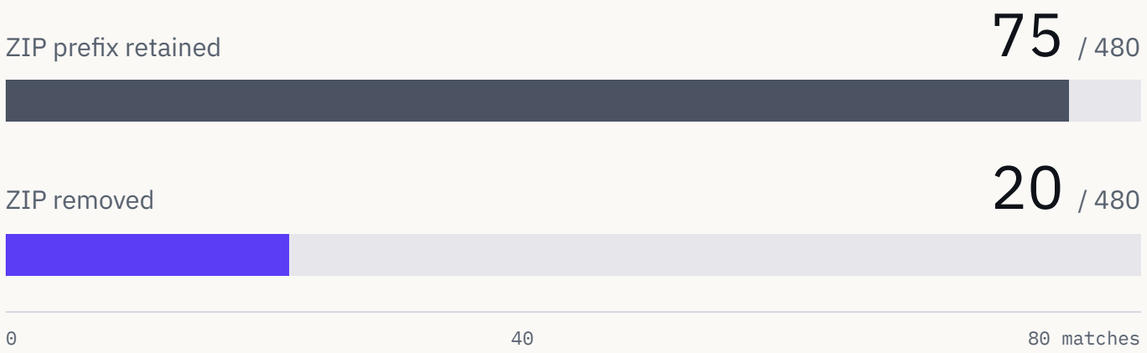

Known membership. Public deterministic QI rules.
Private owner correspondence scores matches after the attack.

Three protected answers stay identical between these two recipes

Six-month failure rate	36.25%
Liver-subgroup failure rate	45.03%
Age–weight correlation	0.6893

All three meet their declared tolerances against the source. Correlation uses interval midpoints.

Baseline: age bands 5 years, ZIP prefix 3, weight bands 5 kg. Modeled source-target matches, not population risk probabilities. Selected answers do not establish full scientific equivalence.

24

configurations requested
in the standard demo search

Actual recipes. Measured tradeoffs.

The deterministic Python engine applies transformations, recalculates disclosure and research metrics, then marks nondominated candidates.

Teams can inspect the recipe, selected answers and failed constraints for every candidate.

01 Apply recipe

02 Measure both objectives

03 Compare the frontier

Finite grammar. Hard cap 80. Small inputs may exhaust available recipes earlier. Frontier status does not imply feasibility or release clearance.

Coverage with clear boundaries

26 threat families

17 catalogued protections

Registry entries disclose implementation and applicability.

Runs today

Public-rule exact linkage, sample uniqueness, sensitive-value inference, research tasks, subgroup diagnostics and bounded recipe search.

Optional pyCANON cross-checks cover categorical k/l/t diagnostics.

Explicitly outside current coverage

Calibrated fuzzy linkage, trained synthetic generation, general membership attacks, reconstruction and persistent differential privacy accounting.

Unavailable analyses abstain. Catalog presence never means assessment success.

React + TypeScript + Vite. Python + FastAPI + pandas + NumPy + SciPy. Archie: public entry + Core PostgreSQL/GraphQL. Vercel: UI assets + Python compute. No generative AI runtime.

A clearer review before teams share data

See what remains identifying.
See what the selected study questions lose.
Carry evidence into internal or partner review.

Public demo

Fictional fixtures only.
Uploads blocked.

Full local version

Private CSVs on your computer.
Protected exports after review.

Human review required. No automatic release clearance, compliance certification or privacy guarantee.

Open the fictional demo on Archie